MVPCompanion

From Idea to MVP: A Worked Example

Follow one imperfect app idea through all 15 stages of progressive clarification.

Case: BookScout → MyBookScout → Matchshelf

A practical companion to the MVPCompanion Method

The central lesson

A good MVP is not simply a long feature list with items removed. It results from clarifying the problem, primary user, desired outcome, product model, Core User Goals, scope, experience, data, requirements, learning goals and technical boundaries.

How to read this example

This resource uses the BookScout → MyBookScout → Matchshelf case documented in the MVPCompanion project. It is a **methodological reconstruction**, not a claim that historical work literally followed the 15 stages in this order. These product names are case artifacts, not method stages.

Evidence discipline: supported case details are distinguished from illustrative hypotheses and unresolved decisions. Do not treat an assumption as validated evidence.

The case at a glance

The starting idea already contains a proposed solution. The method treats it as a hypothesis and asks what must be clarified before building.

Case artifact	Role	What it represents
BookScout	Initial idea	A book app to catalog a library and use it for personalized recommendations.
MyBookScout	Emerging product concept	A broader concept around a reader's library, taste profile, discovery and recommendations.
Matchshelf	Focused MVP	The intended result of defining the smallest coherent product that delivers core value and tests key assumptions.

The 15-stage spine

1 Idea	2 Problem	3 User
4 Value	5 Product	6 Context & Lifecycle
7 Core Jobs	8 Scope	9 Experience
10 Information Architecture	11 Data	12 Requirements
13 Learning	14 Technical Boundaries	15 MVP Boundary

Cross-cutting: Desirability, Differentiation, Viability, Feasibility and Learning; Market Discovery and Business Model; and E/H/D/O labels for Evidence, Hypothesis, Derived conclusion and Open question.

Stage 1 — Idea

Core question: What could this be?

Case setup
“Build an app for book lovers that lets them catalog their books and use their library to receive personalized recommendations.”

Reasoning / illustrative decision
The initial idea combines cataloging and personalized recommendations.

Evidence discipline (E/H/D/O)
H: Readers will value a combined library-and-recommendation product. O: Which part represents the strongest unmet need?

Stage output: Idea Brief: concept, intended audience, benefit and assumptions.

Stage 2 — Problem

Core question: What problem exists independently of the proposed app?

Case setup
The starting idea is solution-oriented, so the problem must be separated from the app concept.

Reasoning / illustrative decision
A plausible problem to investigate is that readers struggle to decide what to read next and find generic recommendations insufficiently personal. This remains a hypothesis.

Evidence discipline (E/H/D/O)
H: Discovery friction matters. O: How often does it occur and what workarounds are used? Gather interviews, observed behavior or other evidence.

Stage output: Problem statement that does not prescribe a catalog or recommender.

Stage 3 — User

Core question: Who experiences the problem most directly?

Case setup
Avoid treating all book lovers as one segment.

Reasoning / illustrative decision
Working hypothesis: an active reader with established preferences who often struggles to find a suitable next book. This segment needs validation.

Evidence discipline (E/H/D/O)
H: This group has a recurring need and motivation to change behavior. O: Which segment has the strongest problem and willingness to try a new approach?

Stage output: Primary User Profile, context, current alternatives, constraints and non-target groups.

Stage 4 — Value

Core question: Why does solving this matter?

Case setup
Translate the problem into an outcome, not a feature list.

Reasoning / illustrative decision
Working value proposition: “Help active readers find their next worthwhile book with recommendations that reflect their taste—not just popularity.”

Evidence discipline (E/H/D/O)
H: Relevance and explanations improve discovery decisions. O: Which benefit matters most: fit, novelty, confidence, convenience or time saved?

Stage output: Value proposition, benefits, alternatives and differentiation hypothesis; Market Discovery and Business Model remain cross-cutting.

Stage 5 — Product

Core question: What is the product fundamentally?

Case setup
Define the product mental model before designing screens.

Reasoning / illustrative decision
MyBookScout treats the reader’s library and taste as foundations for discovery. Conceptual entities in the case include Work, Edition, ISBN, UserBook, Interaction, UserTaste and Recommendation.

Evidence discipline (E/H/D/O)
D: Recommendations require some representation of books, reader preferences and interactions. O: Which entities are essential for the first release?

Stage output: Product Mental Model and high-level entity map; do not confuse it with the MVP data model.

Stage 6 — Context & Lifecycle

Core question: How does the user and product context evolve?

Case setup
Identify meaningful situations, triggers, states and transitions.

Reasoning / illustrative decision
Candidate journey: encounter or add a book, receive a recommendation, decide whether it fits, save or reject it, and possibly provide feedback.

Evidence discipline (E/H/D/O)
H: Feedback may improve relevance. O: Which states and transitions are necessary for the core experience?

Stage output: Context and lifecycle model, limited to states that materially affect the product.

Stage 7 — Core Jobs

Core question: What must users accomplish?

Case setup
Describe outcomes, not interface elements.

Reasoning / illustrative decision
Candidate Core User Goals: represent enough of my reading taste; discover books likely to fit me; understand why a recommendation may fit; save or reject a recommendation.

Evidence discipline (E/H/D/O)
H: These goals form the smallest valuable set. O: Which goal is primary? Validate priorities with target users.

Stage output: Prioritized Core User Goals with trigger, action, desired outcome and minimum information.

Stage 8 — Scope

Core question: What is the smallest useful capability set?

Case setup
Derive capabilities from Core User Goals and challenge every feature.

Reasoning / illustrative decision
Candidate capabilities: add a small starter library; create an initial taste profile; show a limited recommendation list; explain fit; allow Save / Not for me feedback.

Evidence discipline (E/H/D/O)
H: A compact library plus recommendations may create value. O: Is bulk import essential, or can manual selection test the promise first?

Stage output: Feature map classified Must Have / Should Have / Later / Not in MVP, with explicit exclusions.

Stage 9 — Experience

Core question: How does the user reach value end to end?

Case setup
Design a complete journey, not disconnected features.

Reasoning / illustrative decision
Illustrative happy path: start → select familiar books → view a taste summary → receive a short recommendation list with explanations → save or reject a result.

Evidence discipline (E/H/D/O)
H: Users can reach value without lengthy setup. O: How much input is needed? What happens when the profile feels wrong?

Stage output: Core flow, first-value moment, onboarding and essential empty/error states.

Stage 10 — Information Architecture

Core question: How should the product be organized?

Case setup
Navigation should reflect the mental model and core journey.

Reasoning / illustrative decision
The broader concept uses Library → Taste → Discovery → Decision → Learning. A focused MVP should expose only what its selected journey needs.

Evidence discipline (E/H/D/O)
D: If the core journey is library-to-recommendation-to-decision, navigation should make that path obvious. O: Is a guided flow sufficient?

Stage output: Navigation model and screen inventory; avoid exposing every long-term capability.

Stage 11 — Data

Core question: What information must the MVP store or derive?

Case setup
Translate the product model into the minimum implementable data model.

Reasoning / illustrative decision
Potential core entities: book/work reference, selected user book, taste/profile, recommendation and feedback. Rich edition detail or advanced enrichment may be deferred.

Evidence discipline (E/H/D/O)
H: Simplified book data may suffice for early validation. O: Which metadata source is reliable and permitted? What user data must be stored?

Stage output: MVP Data Model: fields, relationships, provenance, ownership and lifecycle.

Stage 12 — Requirements

Core question: What exactly must the system do?

Case setup
Specify observable behavior and testable acceptance criteria.

Reasoning / illustrative decision
Example: “When a user requests recommendations after selecting the required starter books, the system displays a ranked list with a short explanation for each result.”

Evidence discipline (E/H/D/O)
Acceptance example: given adequate input, when recommendations are requested, return a defined-size list with a reason for each result; if input is insufficient, explain the next step.

Stage output: Functional requirements, inputs/outputs, validation, states, dependencies and acceptance criteria.

Stage 13 — Learning

Core question: What must the MVP teach us?

Case setup
Treat the MVP as a learning instrument, not merely working software.

Reasoning / illustrative decision
Candidate hypotheses: users understand the promise; recommendations feel relevant; explanations help decisions; users will provide enough input.

Evidence discipline (E/H/D/O)
Candidate signals: setup completion, recommendation save rate, rejection rate, feedback rate and return use. These are proposed signals, not established benchmarks.

Stage output: Learning questions, hypotheses, experiments, success signals and decision rules.

Stage 14 — Technical Boundaries

Core question: What is technically necessary now?

Case setup
Define enough architecture to estimate and build without solving the long-term system prematurely.

Reasoning / illustrative decision
The broader concept separates metadata, semantic enrichment, retrieval, user-specific scoring, constraint filtering, explanations and feedback. Keep only layers that support the MVP's goals and learning.

Evidence discipline (E/H/D/O)
H: Simpler retrieval and ranking may be enough initially. O: Which metadata services, data rights, privacy and AI components are feasible?

Stage output: Platform, storage, data source, integrations, AI role, risks and technical unknowns.

Stage 15 — MVP Boundary

Core question: What is the smallest coherent product that delivers core value and supports learning?

Case setup
Combine decisions from Stages 7–14 into a defensible boundary.

Reasoning / illustrative decision
Matchshelf is the case name for the focused MVP outcome—not a pre-decided feature list. Its promise, primary user, goals, capabilities, experience, data model, commercial hypothesis and exclusions must be derived.

Evidence discipline (E/H/D/O)
Retain a capability only if it supports a Core User Goal, enables a coherent end-to-end journey or tests a critical assumption. Final exclusions remain decisions to resolve.

Stage output: Integrated MVP Specification and explicit boundary statement. The specification is the result of the 15 stages, not Stage 16.

From clarification to a focused MVP

The result is not simply “MyBookScout with features removed.” The method asks which coherent product expression can deliver selected core value and test important assumptions.

Illustrative end-to-end journey

1 · Start	2 · Seed taste	3 · Discover	4 · Decide
User understands the promise.	Selects or adds familiar books.	Sees a short list with reasons.	Saves or rejects a result; feedback becomes a learning signal.

Illustrative scope boundary

MVP candidate	Possible later	Defer unless validated
Starter-library input; initial taste representation; limited recommendations; explanations; Save / Not for me feedback.	Bulk import, richer library management, extensive personalization controls and advanced analytics.	Social features, comprehensive reading tracker, gamification and complex integrations.

This is a boundary to challenge, not a final Matchshelf specification. Confirm it against evidence, feasibility and learning goals.

Traceability in practice

Problem — Readers struggle to identify a worthwhile next book.
Primary user — A specific active-reader segment (hypothesis).
Outcome — Choose a personally relevant next book with more confidence.
Core User Goal — Discover and evaluate a book that fits my taste.
Capability — Short recommendation list with explanations.
Requirement — Each result includes a reason for its fit.
Acceptance criterion — Each returned result has a non-empty explanation; insufficient input produces a clear next step.
Metric / learning signal — Save rate and user-reported relevance, interpreted with qualitative feedback.

What the example teaches

Clarify before designing

The original idea includes a solution. First investigate the problem, user and desired outcome.

Model before listing features

The product model and Core User Goals help explain which capabilities are necessary.

Scope for coherence, not just smallness

A minimal product still needs a complete path from entry to a meaningful outcome.

Keep uncertainty visible

Assumptions are not evidence. Unknowns should guide research and experiments.

Define learning before building

A working app is not enough. Define what must be learned and what results would change the plan.

Separate method from case history

BookScout, MyBookScout and Matchshelf are case artifacts. The 15 stages are the MVPCompanion Method.

Practical review before handoff

- Can we describe the problem without naming the app or its features?
- Is the primary user specific enough to recruit and learn from?
- Is the desired outcome explicit and meaningful?
- Can every MVP capability be traced to a Core User Goal or learning need?
- Does the experience deliver value end to end?
- Are requirements observable and acceptance criteria testable?
- Are metrics and decision rules defined before the experiment?
- Are technical assumptions, data dependencies and exclusions explicit?
- Are E, H, D and O labels used honestly?
- Could another person understand what to build without repeatedly guessing?

Terminology note

This resource uses the current 15-stage MVPCompanion architecture. Stage 7 is Core Jobs; Stage 11 is Data; Stage 12 is Requirements; Stage 13 is Learning; Stage 14 is Technical Boundaries; Stage 15 is MVP Boundary. The Integrated MVP Specification is the result of the stages, not an additional stage.

Source and status note

Case framing is based on the MVPCompanion Master Case Bible v0.3 and the App Idea to MVP Spec Guide in the project. Unsupported details are intentionally left as illustrative hypotheses or open questions. The proposed MVP scope and metrics are not presented as validated.