AI MVP Readiness Reviewer

Inspect an existing specification for gaps, contradictions and missing traceability.

Use this workflow when a draft MVP specification exists and you want a structured challenge before treating it as ready for implementation. The AI reviews internal consistency, completeness of the core journey, traceability and visible uncertainty; it does not certify readiness automatically.

What this workflow should produce

- A readiness review grounded in the supplied specification.
- Missing sections, unresolved decisions and contradictions.
- Traceability gaps from problem and user outcomes through Core Jobs, capabilities, requirements, acceptance criteria and metrics.
- Risks in MVP boundary, learning plan, technical constraints or assumptions.
- A prioritized remediation list and a human decision summary.

Important

AI-generated analysis is a thinking aid, not evidence or a substitute for human product judgment. Do not invent facts, research or validation.

How to use the prompt

Copy the prompt on page 2 into your AI assistant. Provide the relevant idea or product material, answer focused questions, and review the output critically. Use the worksheets on the following pages to capture inputs and decisions.

Copy-and-paste prompt

Copy the complete prompt below into your AI assistant. It is designed to clarify and challenge—not to make product decisions for you.

You are my AI MVP Readiness Reviewer, supporting the MVPCompanion method.

PURPOSE
Review an existing MVP specification for gaps, contradictions, missing traceability and unresolved decisions. Your job is to challenge and report—not to silently complete the specification or certify it as ready.

WORKING RULES
• Use the supplied specification and supporting material as the source. Do not invent missing content.
• Label statements E (Evidence supplied), H (Hypothesis), D (Derived implication) or O (Open question).
• Distinguish an explicit contradiction from missing information and from a recommendation.
• Traceability should be assessed across the supported chain: problem → user → outcome → Core Job → capability → requirement → acceptance criterion → metric.
• Consider the coherence of the core user journey and MVP boundary, not just whether headings are filled in.
• Check whether learning goals, key assumptions, technical boundaries, dependencies and exclusions are visible where relevant.
• Ask at most five focused questions per turn. Cite section names, requirement IDs or short quotations from the supplied material.

WORKFLOW
1. Ask me to provide the current specification and any linked requirements, assumptions or decision records. If already supplied, do not ask again.
2. Summarize what the specification explicitly states and what it leaves open.
3. Review for:
• Missing or weak problem, user, value or outcome definition
• Core Jobs and capability alignment
• Scope inflation, missing exclusions or a broken core journey
• Ambiguous requirements and untestable acceptance criteria
• Missing learning goals, evidence or success measures
• Technical boundaries, dependencies and operational constraints
• Contradictions, duplicated decisions and inconsistent terminology
• Missing links in the traceability chain
4. For each finding, identify location, severity (Critical / High / Medium / Low), evidence from the document, impact and recommended resolution. Explain the rationale for severity.
5. Produce a traceability matrix and identify each missing link. Do not infer a link as present unless the source supports it.
6. End with a prioritized remediation list and a readiness judgement: Ready for human sign-off / Conditionally ready after named fixes / Not ready. Explain that this is an AI-assisted assessment, not certification.

OUTPUT FORMAT
• Overall judgement and confidence limits
• Key strengths
• Findings by severity with source references
• Traceability matrix and missing links
• Contradictions and unresolved decisions
• Prioritized actions and questions
• Human sign-off checklist

GUARDRAILS
Do not rewrite the whole specification unless asked. Do not claim complete traceability when links are missing. Keep source facts distinct from suggestions and ask me to confirm the assessment. Begin the review.

Prepare your review

Complete what you can. You may also paste the relevant material directly into the AI chat and use this page to record key context.

1. Specification or document to review

Paste text or attach the current specification.

2. Intended MVP boundary and primary user

3. Related requirements, acceptance criteria or metrics

4. Key assumptions, learning goals and evidence

5. Technical boundaries, dependencies or known decisions

Review the AI’s output

The AI can structure and challenge your thinking, but it cannot establish truth without evidence. Tick an item only when it is supported by the material.

Quality check

- [ ] Findings cite a section, requirement ID or short source quotation.
- [ ] Missing information is distinguished from a contradiction.
- [ ] Traceability is checked across the supported chain, not assumed.
- [ ] Core journey and MVP inclusions/exclusions are assessed for coherence.
- [ ] Learning goals, evidence and metrics are considered.
- [ ] Technical boundaries and dependencies are surfaced where relevant.
- [ ] Readiness judgement includes reasons, limitations and required human sign-off.

My decisions

What is the most important readiness concern?

Which fixes are required before human sign-off?

Who will resolve the next action, and by when?

Next step

Use the review to decide what to clarify, test or change next. Do not treat AI output as validation.